



Training



decoders: 8 shared blocks, d = 512, then one head per stream

all panels: held-out sorghum plant *Sorghum_10_04* (data v2): real inputs, the sampled mask, and model outputs at epoch 600; the maize model differs only in its recipe tokens (1 + 28 tokens of 14 floats)

Distillation



init = teacher weights; 80 % of steps distil, 20 % are plain training steps